

Seedling Labs Engineering Intern Craft Case: AI-Powered GitHub Issue Assistant

Welcome to Seedling Labs!

At Seedling, we believe in the immense potential of dreaming, sprouting, and growing. We are a dynamic tech startup specializing in **AI-Native Product Design & Development as a Service (PDaaS)**. While the world is abuzz with evolving AI tools and agents, our core mission is to solve the critical challenge of **productizing and scaling** these innovations for real-world impact. We engineer sophisticated AI tools and agents, effectively "packaging" them into robust capabilities and comprehensive applications. This enables organizations to achieve AI-native transformations not just faster, but *smarter*, significantly reducing their time to market.

Joining Seedling Labs means being at the forefront of AI innovation, tackling complex problems, and seeing your work directly contribute to impactful solutions. We foster an environment of continuous learning, perseverance, and collaboration, where even the smallest beginnings lead to magnificent growth.

Why Join Us?

If you are passionate about building intelligent systems that solve real-world problems, thrive in a fast-paced environment, and want to contribute to a company that's redefining AI product development, Seedling Labs is the place for you. You'll gain invaluable experience by:

- **Innovating with Cutting-Edge AI:** Directly working with and building upon the latest and greatest AI tools, models, and agents.
- **Driving AI-Native Product Development:** Learning to take product ideas from conception to robust, scalable solutions by leveraging an AI-native approach.
- **Mastering Agile & Efficient Delivery:** Developing skills in building high-quality products and solutions rapidly.
- **Engaging in Human + AI Co-creation:** Being a part of a collaborative team that harnesses the power of both human creativity and AI capabilities.

The Craft: AI-Powered GitHub Issue Assistant

Core Philosophy: "Agentic Thinking in a Box"

This craft is designed to be a microcosm of the actual work you'd do at Seedling. It's not just about building a model; it's about creating a small, useful AI agent that integrates different parts (LLMs, data, APIs) to solve a concrete business problem.

Problem Statement:

At Seedling Labs, we move fast. To maintain quality and speed in our development workflow, efficiently understanding and prioritizing new GitHub issues is crucial.

Your task is to build a simple web application that takes a GitHub repository URL and an issue number as input. The application should then use AI to analyze the issue and provide a structured summary.

Core Requirements:

1. **Input UI:** A simple user interface with fields for a public GitHub repository URL (e.g., <https://github.com/facebook/react>) and an Issue Number.
2. **Backend API:** A lightweight API endpoint (e.g., using FastAPI or Flask in Python) that triggers the analysis.
3. **AI Core:** The backend should:
 - Use the GitHub API to fetch the title, body, and comments of the specified issue.
 - Utilize a Large Language Model (LLM) to process this information.
 - Generate a structured output in a specific JSON format.
4. **Output Display:** The UI should display the generated analysis in a clean, readable format.

Required AI-Generated JSON Output Format:

```
{  
  "summary": "A one-sentence summary of the user's problem or request.",  
  "type": "Classify the issue as one of the following: bug, feature_request, documentation, question, or other.",  
  "priority_score": "A score from 1 (low) to 5 (critical), with a brief justification for the score.",  
  "suggested_labels": ["An array of 2-3 relevant GitHub labels (e.g., 'bug', 'UI', 'login-flow')."],  
  "potential_impact": "A brief sentence on the potential impact on users if the issue is a bug."  
}
```

Tech Stack Suggestions (Flexibility Encouraged):

- **Backend:** Python (FastAPI, Flask)
- **Frontend:** Streamlit or a simple HTML/JS page that calls the API. (Streamlit is excellent for rapid prototyping).
- **LLM Interaction:** LangChain, Llamaindex, or direct API calls (e.g., Google Gemini, OpenAI, or an open-source model via Hugging Face).

Submission:

The entire project should be submitted as a link to a public GitHub repository. Your repository **must** include a clear and comprehensive `README.md` file explaining how to set up and run your project in under 5 minutes. This demonstrates professionalism and empathy for other developers.

Evaluation Framework (The Rubric)

We will evaluate your submission based on the following criteria, reflecting Seedling Labs' core values of problem-solving, quality, speed, and communication:

1. **Problem Solving & AI Acumen (40%)**
 - **Prompt Engineering:** How effectively did you craft the prompt for the LLM? Is it robust and does it reliably produce the desired JSON format? Did you consider techniques like few-shot prompting?
 - **System Design:** How did you structure your code to fetch data, process it with the LLM, and return the result? Is the design logical and efficient?
 - **Handling Edge Cases:** Did you consider scenarios like issues with no comments, or very long issue bodies? How did your solution account for these?
2. **Code Quality & Engineering Practices (30%)**
 - **Clarity & Readability:** Is the code clean, well-commented, and easy to understand?
 - **Project Structure:** Is the project organized logically into files and folders?
 - **README:** Is your `README.md` clear and comprehensive, allowing for easy setup?
 - **Dependency Management:** Did you use a `requirements.txt` or `pyproject.toml` file?
3. **Speed & Efficiency (20%)**
 - **Tool Usage:** Did you leverage libraries (e.g., LangChain, FastAPI) effectively to avoid reinventing the wheel and accelerate development?
 - **Functionality:** Is the final product snappy, usable, and does it fully address the problem statement?
4. **Communication & Initiative (10%)**
 - **Git History:** Are your commit messages clear, descriptive, and do they tell a story of your development process?
 - **Going the Extra Mile:** Did you add any small, thoughtful features not explicitly asked for (e.g., a button to copy the JSON, basic error pop-ups in the UI, caching results)?

Good luck!